

Ser**IoT**ics

Serialization metrics in an
Internet of Things environment

SHAUN DONACHY


```
{ "minType" : "BATHROOM",  
  "majType" : "HUMANITARIAN",  
  "major_area_num" : 19,  
  "minor_area_num" : 436,  
  "quantity" : 3.7826,  
  "unique_id" : 1000,  
  "runtime" : 1443200834318,  
  "item_sensed" : "toilet",  
  "subject_measured" : "water",  
  "sensor_location_name" : "bathroom",  
  "ticks_since_turn_on" : 2 }
```



```
{ "minType": "BATHROOM",  
  "majType": "HUMANITARIAN",  
  "major_area_num": 19,  
  "minor_area_num": 436,  
  "quantity": 3.7826,  
  "unique_id": 1000,  
  "runtime": 1443200834318,  
  "item_sensed": "toilet",  
  "subject_measured": "water",  
  "sensor_location_name": "bathroom",  
  "ticks_since_turn_on": 2 }
```

**JSON message
~ 265 bytes**


```
{ "minType": "BATHROOM",  
  "majType": "HUMANITARIAN",  
  "major_area_num": 19,  
  "minor_area_num": 436,  
  "quantity": 3.7826,  
  "unique_id": 1000,  
  "runtime": 1443200834318,  
  "item_sensed": "toilet",  
  "subject_measured": "water",  
  "sensor_location_name": "bathroom",  
  "ticks_since_turn_on": 2 }
```

JSON message
~ 265 bytes

ProtoBuf Byte Array
~ 69 bytes

Avro Byte Array
~ 58 bytes


```
{ "minType": "BATHROOM",  
  "majType": "HUMANITARIAN",  
  "major_area_num": 19,  
  "minor_area_num": 436,  
  "quantity": 3.7826,  
  "unique_id": 1000,  
  "runtime": 1443200834318,  
  "item_sensed": "toilet",  
  "subject_measured": "water",  
  "sensor_location_name": "bathroom",  
  "ticks_since_turn_on": 2 }
```

JSON message
~ 265 bytes

ProtoBuf Byte Array
~ 69 bytes


Avro Byte Array
~ 58 bytes

Avro Byte Array with schema header
~ 540 bytes

{ Schema }



{ Schema }



Producer Serialization

{ Schema }



Java™


{ Schema }



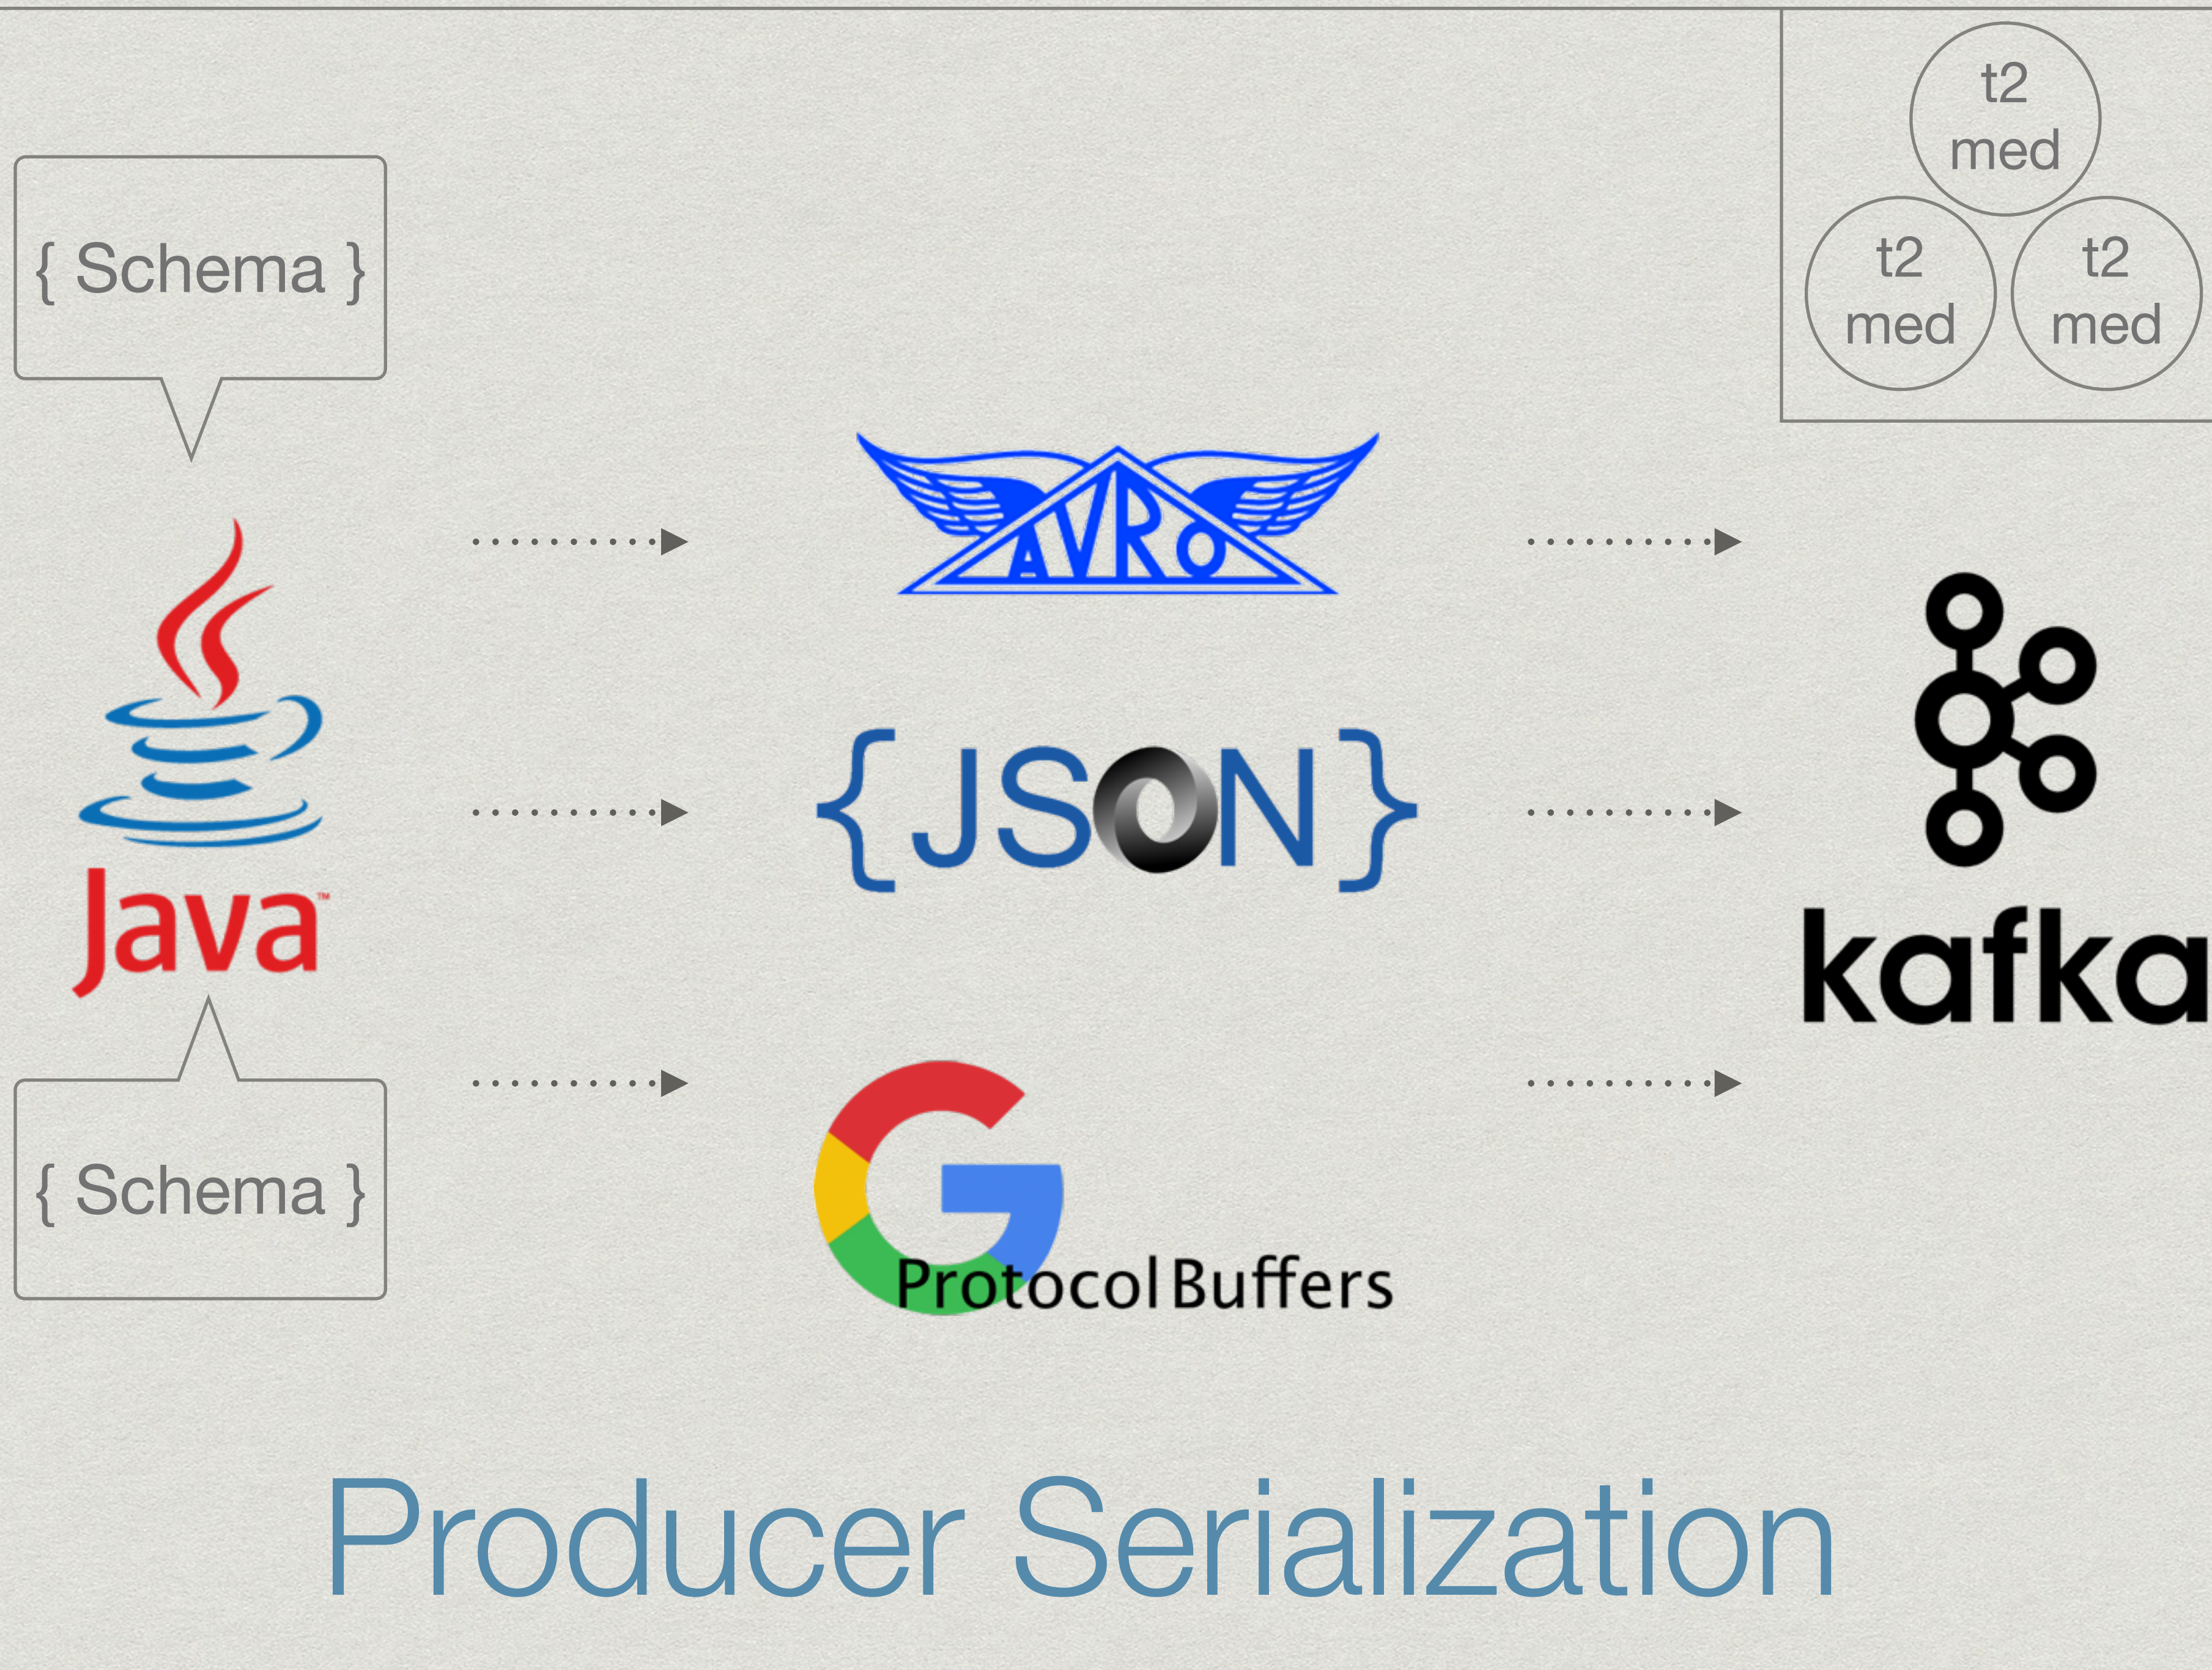
{JSON}




Protocol Buffers



Producer Serialization







Consumer Deserialization






DStream
Micro Batches



Consumer Deserialization






DStream
Micro Batches



Consumer Deserialization



DStream
Micro Batches



Consumer Deserialization

COMPARISON OF SERIALIZATION AND BATCH SIZE

3 instances, 3 JVMs per instance



Challenges

- * Serializing in Java and de-serializing in Python - no docs
- * Java kafka producer API is new - conflicting docs
- * Computing throughput
- * Debugging in Pyspark
- * Parallel reads with receiver-based consumer



SHAUN DONACHY


MS COMPUTER SCIENCE, VIRGINIA COMMONWEALTH UNIVERSITY



VCU

VIRGINIA COMMONWEALTH UNIVERSITY

Spark
Streaming



 RethinkDB



JavaScript



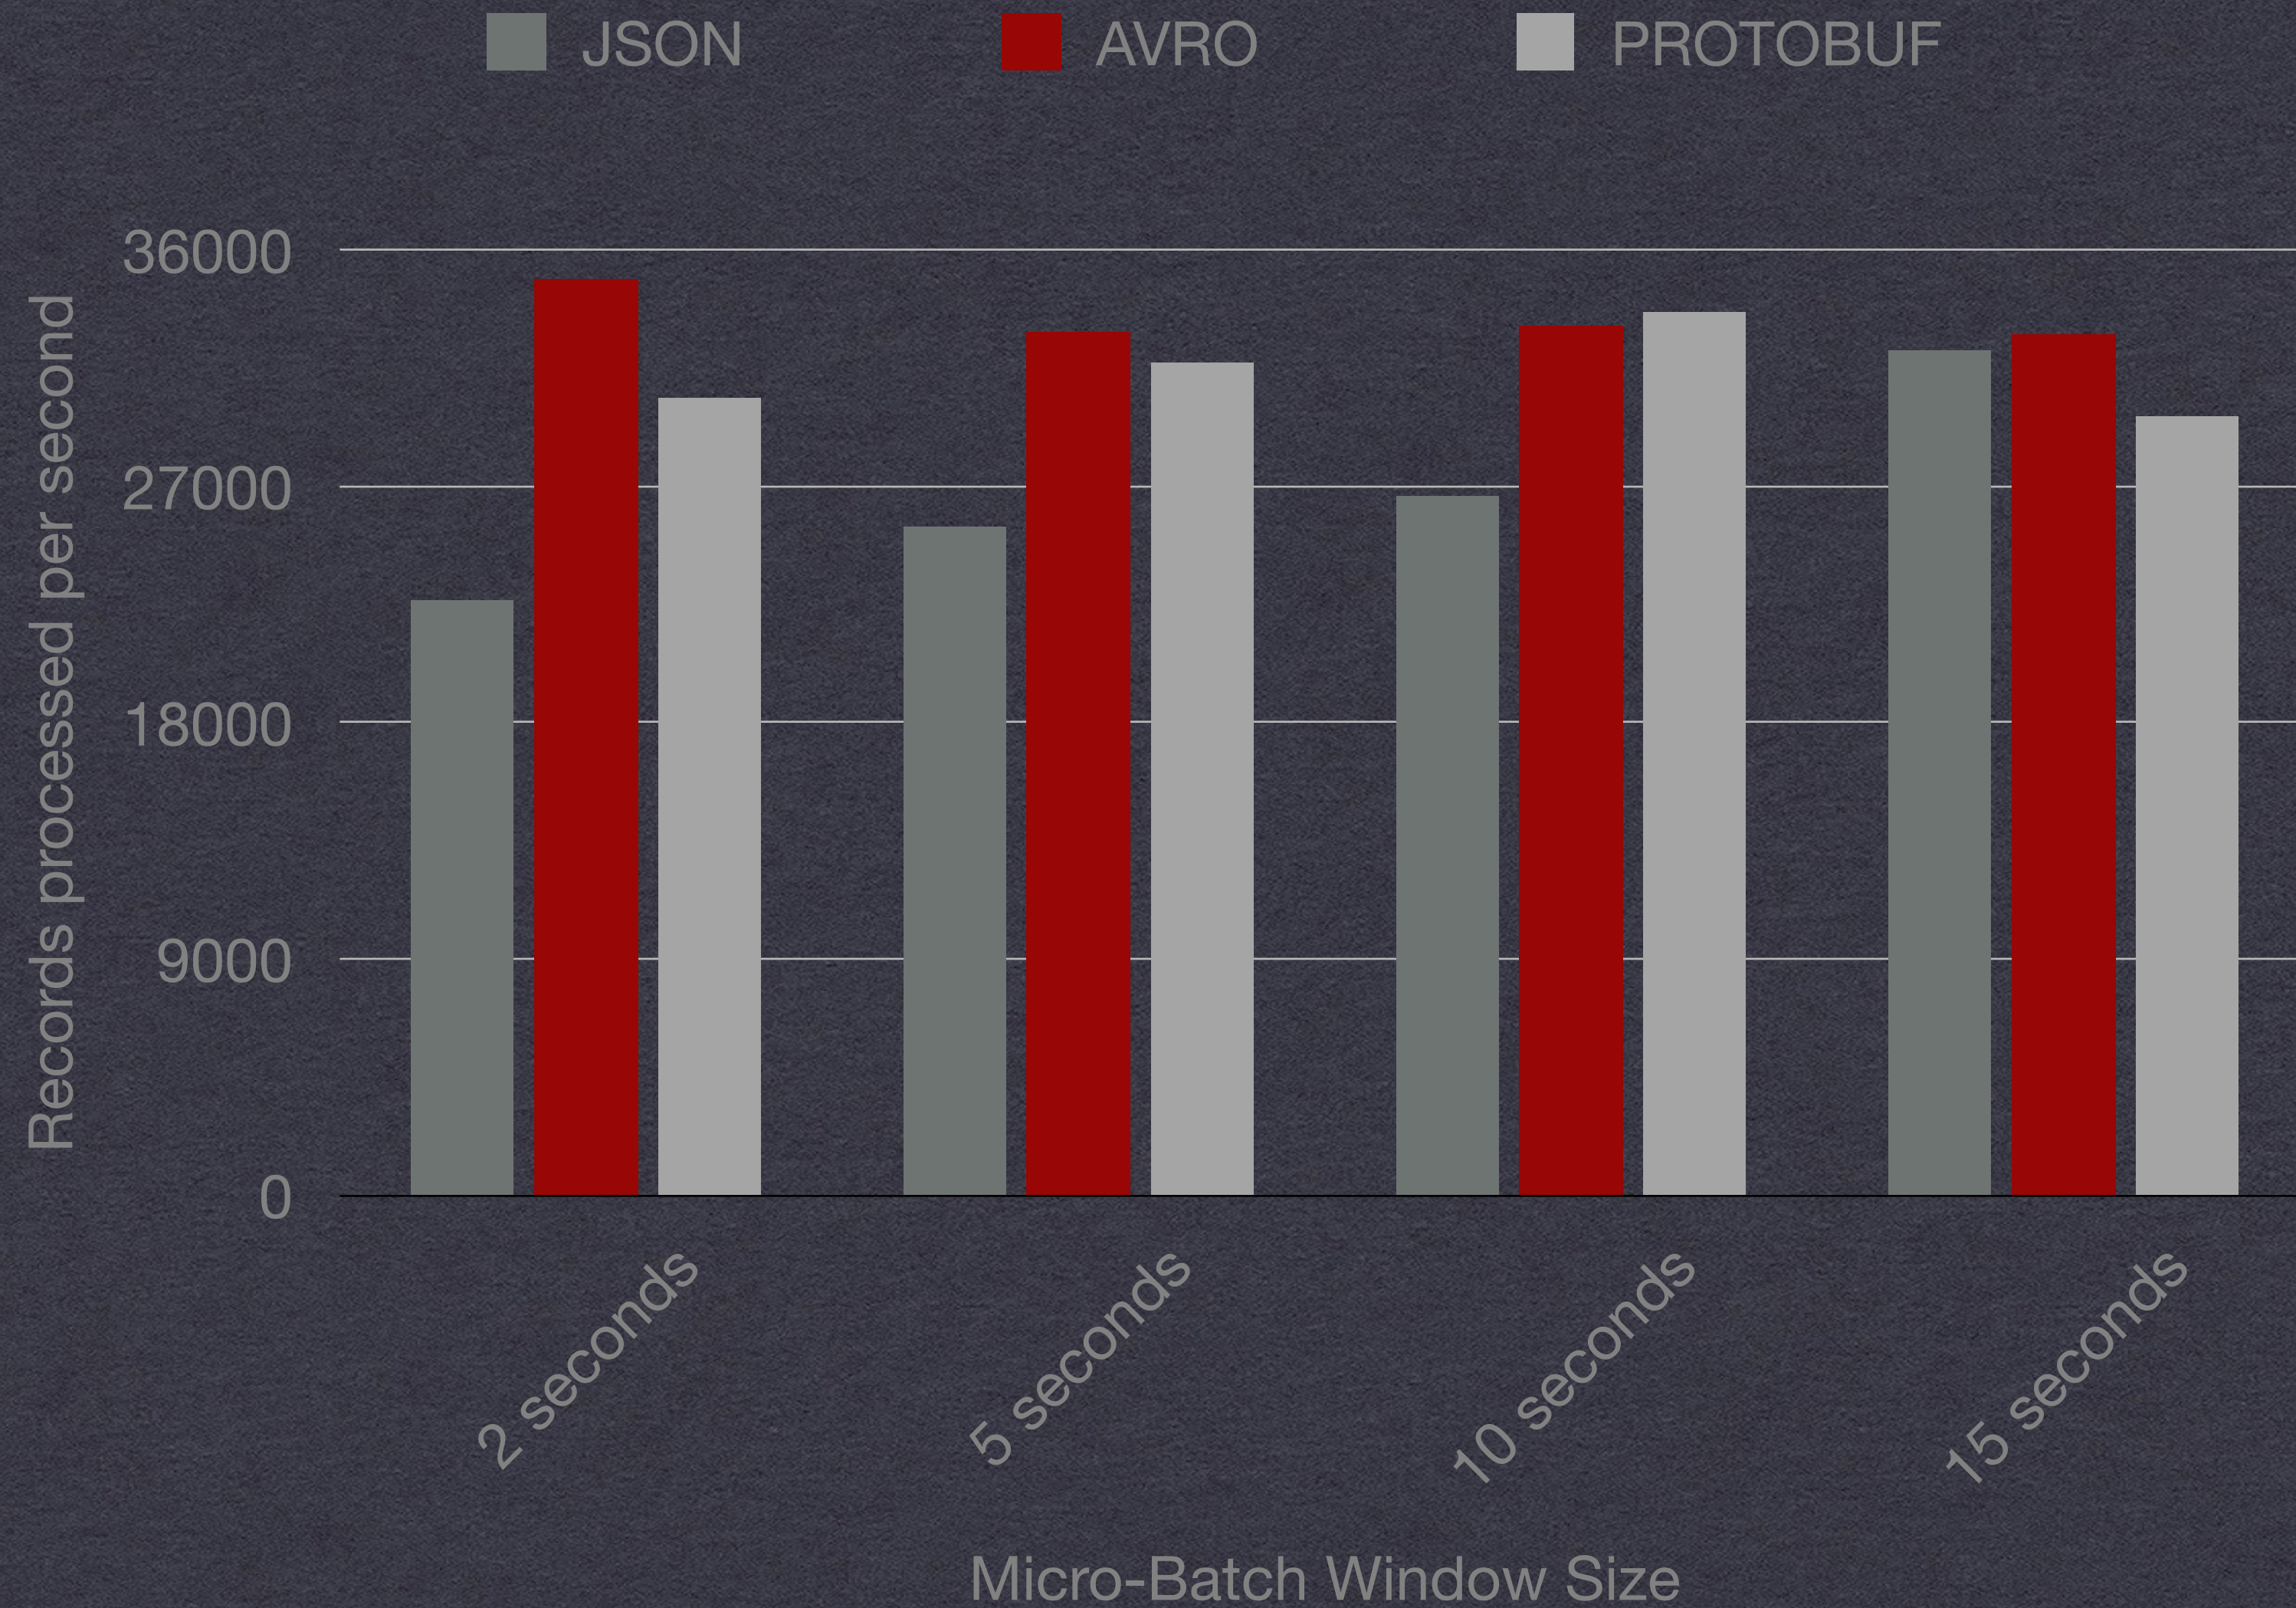
Experimental Parameters

Producer Machines	Producer jvm per machine	Thread spawn rate	Total Sensors per jvm
2/3	1/3	5 sec	125,000

Kafka Partitions/Replication	Spark Read Parallelism	Spark Compute Parallelism	Spark micro-batch window
6/2	6	6	2/5/10/15 sec

COMPARISON OF SERIALIZATION AND BATCH SIZE

2 instances, 1 JVM per instance



Findings

- * Avro and Protobuf serialization greatly improves throughput over JSON
- * Overall Avro and Protobuf show similar performance, the differences are in their characteristics

Image credits

- * Disney, Star Wars
- * [BrainlessTales.com](#)
- * Settlers of Catan
- * Apache: Spark, Avro, Kafka
- * Python, rethinkDB, flask